

GT Assure

AI-assisted control testing and PBC management for an assurance practice. Everything runs on the consultant's machine; no client evidence leaves it.

The system compresses the mechanical part of control testing — retrieval, first-pass assessment, citation, PBC administration — and leaves the judgement with a named human who approves every conclusion before it can be exported.




Fig. 1 – one folder per client; one network call in the system.

Constraints, not aspirations

GUARANTEE	HOW IT HOLDS
Fully local	All inference via Ollama on the machine. No external call anywhere in the data path. Works offline after setup.
Framework-agnostic	NIS2, DORA, ISO 27001, DigiD, client-specific — loaded as configuration files, never hard-coded.
Human in the loop	The system drafts, a named consultant approves, nothing exports unapproved. It never contacts a client.
Auditable	Append-only log of every analysis: model + version, inputs, evidence consulted, output hash, reviewer.
Pinned model	One generation tag per run. Config records the expected digest; startup verifies and warns on mismatch.
Synthetic data only	Ships with a full fictional engagement, so the demo never needs real client data.




Fig. 2 – the per-control pipeline. The empty branch never reaches the model.

Two steps exist because of a specific failure. Retrieval returns *whole documents* in page order rather than the closest paragraph, because the incident plan's severity matrix and its named crisis roles sit on different pages — judging the control on whichever page embedded closest manufactures false gaps. And if retrieval returns nothing, the second model call never happens, because a model with no evidence in front of it will answer from general knowledge about what such a control usually looks like.

Sample testing: who does what

STEP	WHO
Sample size	a lookup in the firm's table (<code>sampling.yaml</code>)
Which items are drawn	seeded draw, reproducible from engagement + control + population hash
Which attributes exist	the framework file
Evidence corroborates an item	the local model — one call per item
Exceptions, conclusion, workpaper	code

Three interlocks, each written after watching an 8B do the thing it prevents

INTERLOCK	THE FAILURE
Only passages that literally name the item reach the model	On the first live run it concluded all twelve items established, citing September release documentation for a February change.
A conclusion without a resolvable citation is downgraded	Fluent agreement with no traceable source reads identically to a real conclusion in a filed workpaper.
n.v.t. refused where the client's export asserts the attribute	"Not applicable" is how real gaps quietly stay out of the exception count.

Evidence is found *literally*, not semantically. "Which passages mention ticket 14721" is a text question; a ticket number carries almost no embedding signal, so the correct document repeatedly fell outside top-k while unrelated ones ranked above it.

Measured on qwen3:8b

BEHAVIOUR	ACROSS TWO RUNS
Changes absent from the release documentation	Stable — <i>niet vastgesteld</i> on every attribute, both runs.
CAB approval, test environment	Stable.
Acceptance attribute on clean items	Not stable — four of eight flagged as exceptions in one run that should have passed.
A control with no evidence at all	The 14B reliably says <i>insufficient evidence</i> ; the 8B is more likely to accept an adjacent document.

Errors run conservative: the tool over-refers to the reviewer rather than passing items silently. That is the right direction for assurance work, and it is not accuracy.

Pitch the cited reasoning, not the label. That is the demo instruction and it is also the honest description of what an 8B is good for.

Walkthroughs

A transcript becomes the process description and a flowchart. The transcript is numbered by line and **every extracted step must cite the line it came from**; a step citing nothing is discarded and counted. Exact rather than fuzzy on purpose — a line number either resolves to something somebody said or it does not. Control references are whitelisted against the control's own identifiers, steps are renumbered by position because the model's numbering drifts, and the flowchart is laid out deterministically so two workpapers for the same process can be compared.

Push on this

- Approval is enforced by writing the export file only at approval time. Elegant, but the guarantee lives in a file's existence rather than in a check I can point at.
- The 8B's attribute-level instability is measured but not bounded. Unknown on a real population an order of magnitude larger.
- The audit chain is append-only JSONL on the same disk as the evidence. Tamper-evident, not tamper-proof.