

GT Proposal

Drafting proposals from the ones we already wrote, and pricing them from what comparable engagements actually cost. Locally, with no proposal, client name or fee leaving the machine.

A proposal takes 40–100 hours and most are lost, which makes the firm's pile of past proposals both its largest body of reusable writing and its largest pile of sunk cost.




Fig. 1 – the model proposes; the code decides.

The estimate

STEP	WHO
Which past engagements are comparable	the model — selection only
Why each is comparable	the model, one clause each
Median and range of hours and fee	code, from recorded actuals
Implied hourly rate	code
Win rate among the comparables	code
What the lost bids were priced at	code

The model is **never shown a fee or an hour count** when it selects — only file, client, service line and scope tags — so it judges similarity of work rather than anchoring on someone else's price. It replies with indices into a numbered list, and anything out of range is discarded silently, because an out-of-range index is the model inventing an engagement.

A model asked to "estimate the fee" returns a plausible number with no provenance. This returns one you can trace to named engagements, recompute by hand, and defend in a pricing discussion.

Measured on qwen3:8b

Selection took ~10 s and returned the two engagements a person would have picked. On a re-run it also selected a third — reasonably, the only utility in the corpus — which moved the median. **The comparable set is not stable across runs.** So the named comparables are the output and the median is a summary of them, not the other way round.

Retrieval is restricted to the comparables

A probe like *"deliverables, milestones, indicative timeline"* matches every proposal's deliverables table about equally well. The first live run drafted a water authority's internal audit timetable as a payments company's ISAE 3402 readiness plan — table intact, week numbers intact, "Readiness confirmation memorandum, Week 15" and all. Drafting now retrieves only from the comparables already selected for the estimate.




Fig. 2 – three code checks, none of which trust the model to have followed instructions.

CHECK	THE FAILURE IT WAS WRITTEN FOR
Confidentiality	Full name, name without legal form, and the one distinctive word people actually say. The third form is the one that matters: the check initially missed <i>"Caldbeck manages approximately 6,400 units across two regions"</i> , lifted whole, because the catalogue's name is "Caldbeck Housing Association".
Attribution	A section drafted from six extracts that cites none of them. Fluent, original-reading, entirely-assembled prose is the failure mode of a reuse tool.
Foreign specifics	The name check only caught a borrowed <i>name</i> ; a borrowed figure with the name stripped out passed it.

A draft with an outstanding name **cannot be approved** — the call refuses rather than warning. What it still cannot catch is a paraphrased specific: "around six and a half thousand units" for "6,400 units" passes, because matching is exact. That is deliberate — a similarity threshold is a dial for how often a confidentiality check fails quietly.

Qualification, before any writing

Entirely code, computed from recorded outcomes. It scores the *firm's* position, not the client.

SIGNAL	COMPUTED FROM
Track record in the service line	won / lost among all catalogued bids in that line
Depth of comparable work	how many selected comparables were won
Recency	months since the most recent comparable engagement
Price pressure	comparables whose loss reason names price
Capability coverage	scope the brief mentions that no comparable covers

Two weak signals is a no-go; one is go-with-conditions. Every signal names the engagements it was computed from, because a partner has to be able to argue with it. Capacity, relationship, independence and strategic value are deliberately *not* scored and print as the human's — the tool cannot see a calendar or a conflicts register.

Push on this

- Hours, fee and outcome are entered by hand, once, per proposal. The whole estimate rests on that being honest and on recording actuals rather than quotes.
- Citation coverage swings between one and four of five sections run to run.
- Exact matching in the confidentiality check is deliberate and arguable. I would like to be talked into or out of it with something better than "add a threshold".