

GT Proposal

Drafting proposals from the ones we already wrote, and pricing them from what comparable engagements actually cost. Locally, with no proposal, client name or fee leaving the machine.

A proposal takes 40 to 100 hours to produce and most of them are lost, which makes the firm's pile of past proposals simultaneously its largest body of reusable writing and its largest pile of sunk cost. GT Proposal indexes that pile and does three things with it: drafts a new proposal section by section from the language of bids that were won, prices it from the recorded actuals of comparable engagements, and refuses to let another client's name out of the building.




Fig. 1 – the model proposes; the code decides.

The estimate is the defensible part

The model does exactly one job in the pricing path: it says which past engagements are comparable, and gives one clause of reasoning each. Every hour, euro and percentage — median, range, implied rate, win rate among the comparables, what the lost bids were priced at — is computed by code from the catalogue's recorded actuals.

Two details make that hold. The model is **never shown a fee or an hour count** when it selects, only file, client, service line and scope tags, so it judges similarity of work rather than anchoring on someone else's price. And it replies with indices into a numbered list: anything out of range is discarded silently, because an out-of-

range index is the model inventing an engagement, and an invented comparable would feed a real number into a real price.

A model asked to "estimate the fee" returns a plausible number with no provenance. This returns a number you can trace to named engagements, recompute by hand, and defend in a pricing discussion.

Measured, on qwen3:8b

On a first-time internal audit brief for a regional water authority, selection took about ten seconds and returned the two engagements a person would have picked. On a re-run it also selected a third — reasonably, it is the only utility in the corpus — which moved the median. **The comparable set is not stable across runs.** So the named comparables are the output and the median is a summary of them, not the other way round.

Retrieval is restricted, and that is what makes a section right

This one was learned the hard way. A probe like *"deliverables, milestones, indicative timeline"* matches every proposal's deliverables table about equally well, so the first live run drafted a water authority's internal audit timetable as a payments company's ISAE 3402 readiness plan — table intact, week numbers intact, "Readiness confirmation memorandum, Week 15" and all. Drafting now retrieves only from the comparables already selected for the estimate, and widens to the whole corpus only when that filter returns nothing.




Fig. 2 – three code checks, none of which trust the model to have followed instructions.

The interlocks, and why each one exists

CHECK	THE FAILURE IT WAS WRITTEN FOR
Confidentiality	Full name, name without legal form, and the one distinctive word people actually say. The third form is the one that matters: the check initially missed <i>"Caldbeck manages approximately 6,400 units across two regions"</i> , lifted whole from another engagement, because the catalogue's name is "Caldbeck Housing Association".
Attribution	A section drafted from six extracts that cites none of them. Fluent, original-reading, entirely-assembled prose is the failure mode of a reuse tool, and that is exactly what an uncited section looks like.
Foreign specifics	The name check only ever caught a borrowed <i>name</i> ; a borrowed figure with the name stripped out passed it. Every number and name-shaped phrase must appear in the brief or in an extract that section actually cites.

A draft with an outstanding name **cannot be approved** — the approval call refuses rather than warning. Citation tags are stripped before the confidentiality check runs, because a tag contains the past client's name and, left in, every correctly cited sentence flagged itself.

What it still cannot catch is a *paraphrased* specific: "around six and a half thousand units" for "6,400 units" passes, because matching is exact. That is the same deliberate choice the name check makes — a similarity threshold is a dial for how often a confidentiality check fails quietly. Read every number in a draft against the brief. That is the reviewer's job and the tool does not remove it.

Before the writing: qualification

The most valuable thing this does is not write proposals faster. It is to stop ten of them going into a bid the firm was never going to win. Qualification is **entirely code**, computed from recorded outcomes, and it scores the firm's position rather than the client: track record in the service line, depth of comparable work, recency, price pressure, and capability coverage. Two weak signals is a no-go; one is go-with-conditions. Every signal names the engagements it was computed from, because a partner has to be able to argue with it.

Four things are deliberately not scored and are printed as the human's: capacity, relationship, independence and conflicts, and strategic value. The tool cannot see a calendar or a conflicts register, and a score that silently omitted them would read as though it had covered them.

What I would want a reviewer to push on

- Hours, fee and outcome are entered by hand, once, per proposal. The whole estimate rests on that data entry being honest and on recording actuals rather than quotes.
- Citation coverage swings between one and four of five sections run to run. The attribution check exists because of that, but the underlying instability is unaddressed.
- Exact matching in the confidentiality check is a deliberate, arguable choice. I would like to be talked into or out of it with something better than "add a threshold".